



GITOPS FOR PLATFORM ENGINEERING

# GitOps Architecture, Patterns and Anti-Patterns

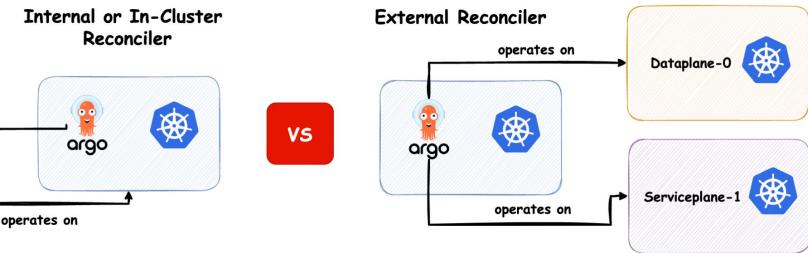
MODULE 03



# GitOps Architecture

# GitOps Architecture

Remember this.



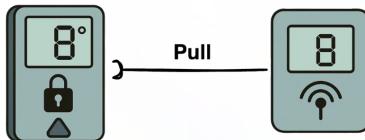
## Declarative



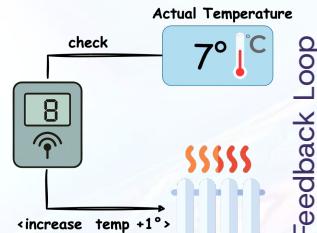
## Versioned and Immutable



## Pull-Based

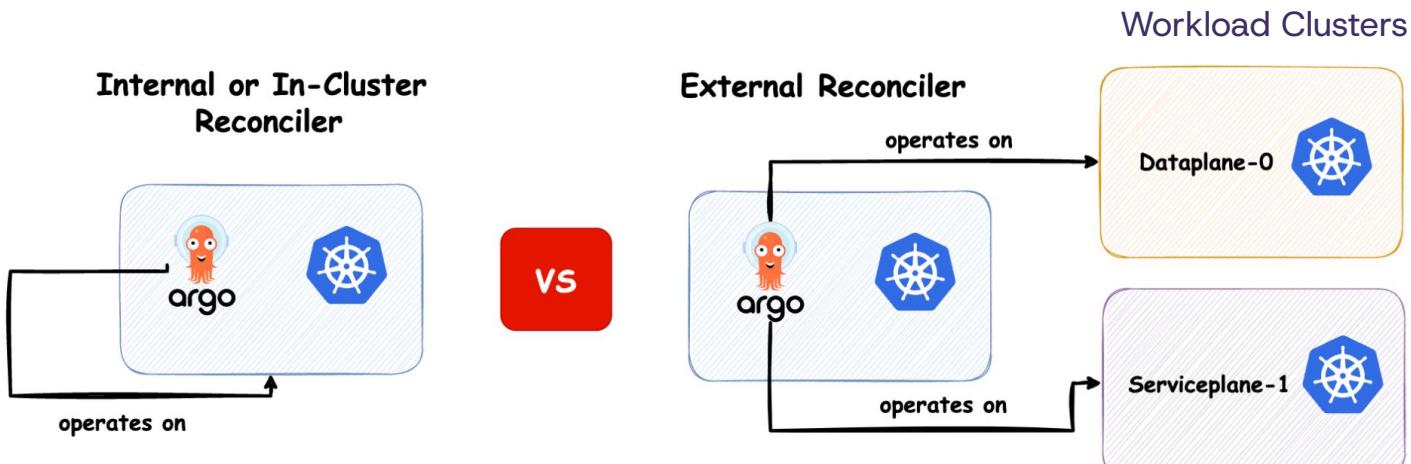


## Continuously Reconciled

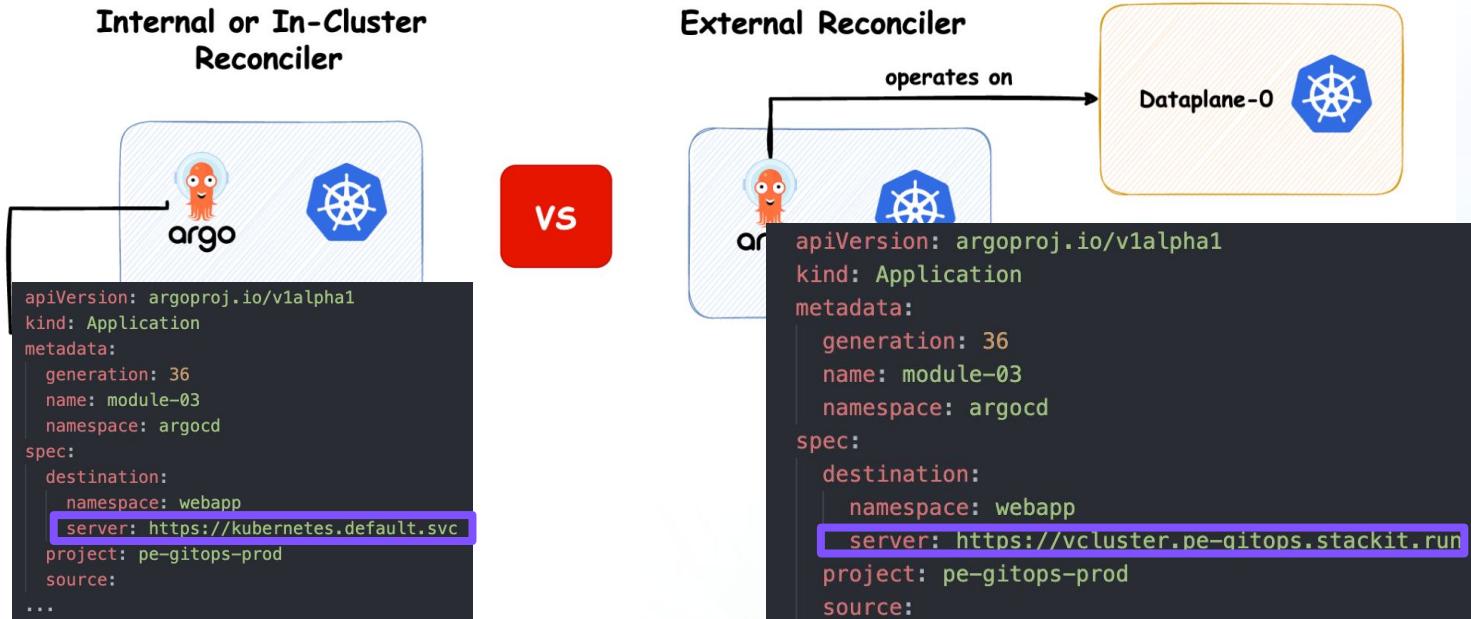


Just four simple principles—no more, no less.

# Internal vs External Reconciler



# Internal vs External Reconciler



# Declarative

Declarative means defining the desired "what" (the end state) and letting the system handle the "how" (the steps to get there).

Definition (YAML): `replicas: 3`

Imperative is a list of commands; Declarative is a description of the destination

`kubectl scale deployment my-app --replicas=3`

Imperative is like giving a driver turn-by-turn directions;  
Declarative is giving the driver an address and letting them find the best route

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
  namespace: webapp
spec:
  replicas: 3
  selector:
    matchLabels:
      app: frontend
  template:
    metadata:
      labels:
        app: frontend
    spec:
      containers:
      - name: frontend
        image: ghcr.io/stefanprodan/podinfo:6.9.4
        ports:
        - containerPort: 9898
      resources:
        requests:
          cpu: 100m
          memory: 32Mi
        limits:
          cpu: 1000m
          memory: 128Mi
```



# Versioned and Immutable

Versioned and immutable means **every change is identified by a unique, permanent tag** (like a Git hash or Docker digest) so that once a version is deployed, it can never be altered—only replaced by a newer version.

## ✗ The Bad Way (Mutable)

```
- name: frontend
  # DANGEROUS: You don't know what's inside
  image: ghcr.io/stefanprodan/podinfo:latest
```

## ✓ The Good Way (Versioned & Immutable)

```
- name: frontend
  # Versioned: Clear release tag
  image: ghcr.io/stefanprodan/podinfo:6.9.4
  # OR even better: Immutable via Digest (SHA)
  # image: ghcr.io/stefanprodan/podinfo@sha256:82138e053dc84...
```



Tags are for humans (versioned); Hashes are for machines (immutable); Never use 'latest' in production.

# Pull-Based

In a Pull-based system, an agent inside the cluster continuously monitors Git and 'pulls' changes to synchronize the state, ensuring the cluster always matches the source of truth.



Push-based is 'Fire and Forget' from the outside; Pull-based is 'Watch and Sync' from the inside.

## ✗ The Push Model (Traditional CI/CD)

- **Code:** `git push` -> CI/CD Server runs `kubectl apply -f deployment.yaml`
- **Risk:** You have to store sensitive Kubernetes credentials (Kubeconfig) inside your CI tool.

## ✓ The Pull Model (GitOps)

- **Code:** `git push` -> Argo CD (inside the cluster) detects the change and pulls it.
- **Benefit:** No cluster credentials leave the cluster. Even if someone manually deletes a pod, the Pull-based agent will "pull" the correct state back.



# Continuous Reconciled

Continuously Reconciled means the system constantly compares the observed state of the cluster with the desired state in Git and automatically fixes any discrepancies (drift) without human intervention.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: argocd-cm
  namespace: argocd
  labels:
    app.kubernetes.io/name: argocd-cm
    app.kubernetes.io/part-of: argocd
data:
  timeout.reconciliation: 180s #3 min
```

```
apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: frontend-app
spec:
  syncPolicy:
    automated: # This enables Continuous Reconciliation
      prune: true    # Automatically delete resources removed from Git
      selfHeal: true # Automatically fix manual changes in the cluster (Drift)
```

## The "Loop" Logic

- ◆ **Observe:** Argo CD looks at the running Pods in your `webapp` namespace.
- ◆ **Diff:** Argo CD looks at the YAML in your Git repo.
- ◆ **Act:** If a Pod was manually deleted or a replica count was changed via `kubectl`, Argo CD "reconciles" it by re-applying the Git state.



Continuous Reconciliation turns 'I hope it's running' into 'I know it's running' by constantly killing the drift.

Context: ske-pe-pro [K8s]  
Cluster: ske-pe-pro  
User: ske-pe-pro  
K9s Rev: v0.50.6 ⚡ v0.50.16  
K8s Rev: v1.34.2  
CPU: 3%  
MEM: 43%

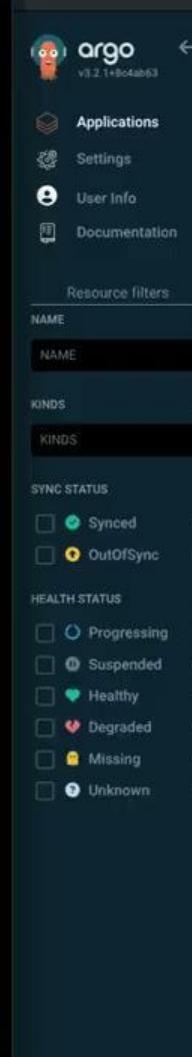
<0> all  
<1> webapp  
<2> external-dns  
<3> kyverno  
<4> mgmt  
<5> projectsveltos



pods(webapp)[3]											
NAME	PF	READY	STATUS	RESTARTS	CPU	MEM	%CPU/R	%CPU/L	%MEM/R	%MEM/L	IP
frontend-b5fd785df-jkd5j	●	1/1	Running	0	2	14	2	0	44	11	100.6
frontend-b5fd785df-xkfw2	●	1/1	Running	0	2	18	2	0	59	14	100.6
frontend-b5fd785df-zm5br	●	1/1	Running	0	2	16	2	0	52	13	100.6

<namespace> <pod>

argo v3.2.1+8c4ab63



Applications

Settings

User Info

Documentation

Resource filters

NAME

KINDS

SYNC STATUS

- Synced 4
- OutOfSync 0

HEALTH STATUS

- Progressing 0
- Suspended 0
- Healthy 6
- Degraded 0
- Missing 0
- Unknown 0

APPLICATION DETAILS TREE

Applications / Q module-03

DETAILS DIFF SYNC SYNC STATUS HISTORY AND ROLLBACK DELETE REFRESH

APP HEALTH APP Sync Status LAST SYNC

Healthy Synced Sync OK

Sync OK to HEAD (SSa9693) Sync OK to HEAD (SSa9693)

Succeeded 4 minutes ago (Tue Dec 23 2025 01:23:46 GMT +0100)  
Author: Artem Lajko <110809333+la-cc@u...

Author: Artem Lajko <110809333+la-cc@u...





# Progressive Delivery



# Progressive Delivery

Is the umbrella term. It's the overarching philosophy or discipline of modern software releases.



## Rolling Update

Replaces pods one by one (incrementally).

**Standard availability** with minimal resource overhead.



## Shadow/Mirroring

Mirrors live traffic to the new version without the user seeing the result.

**Performance & Load** testing under real conditions.



## A/B Testing

Distributes traffic based on user attributes (region, browser, ID).

**Business metrics** (Which version converts better?).



## Blue/Green

Two identical environments; switches 100% traffic from Blue (old) to Green (new).

**Zero Downtime** and instant rollbacks.



## Canary

Routes a small percentage of traffic (e.g., 5%) to the new version to test it.

**Risk mitigation** by testing on **real users**.



While "Continuous Delivery" focuses on moving code from the developer to production as fast as possible, Progressive Delivery focuses on how that code is exposed to users to minimize risk

# Progressive Delivery

Is the umbrella term. It's the overarching philosophy or discipline of modern software releases.



## Rolling Update

Replaces pods one by one (incrementally).

**Standard availability** with minimal resource overhead.



## Shadow/Mirroring

Mirrors live traffic to the new version without the user seeing the result.

**Performance & Load** testing under real conditions.



## A/B Testing

Distributes traffic based on user attributes (region, browser, ID).

**Business metrics** (Which version converts better?).



## Blue/Green

Two identical environments; switches 100% traffic from Blue (old) to Green (new).

**Zero Downtime** and instant rollbacks.



## Canary

Routes a small percentage of traffic (e.g., 5%) to the new version to test it.

**Risk mitigation** by testing on **real users**.

→ While "Continuous Delivery" focuses on moving code from the developer to production as fast as possible, **Progressive Delivery** focuses on how that code is exposed to users to minimize risk



# Trunk-Based vs Branch-Based Development



# Trunk-Based vs Branch-Based Development

## Trunk-Based Development

In this model, developers collaborate on a single branch (usually `main` or `master`). Short-lived feature branches are merged as quickly as possible.

- ◆ **How it works in GitOps:** Argo CD tracks the `main` branch. As soon as a commit is pushed or merged into `main`, Argo CD automatically syncs those changes to the cluster.
- ◆ **Focus:** Speed, Continuous Integration, and avoiding "merge hell."
- ◆ **Best for:** Teams with high automation, automated testing, and a "fail forward" mentality.

## Branch-Based Development (Env per Branch):

In this model, different Git branches represent different environments (e.g., a `staging` branch and a `production` branch).

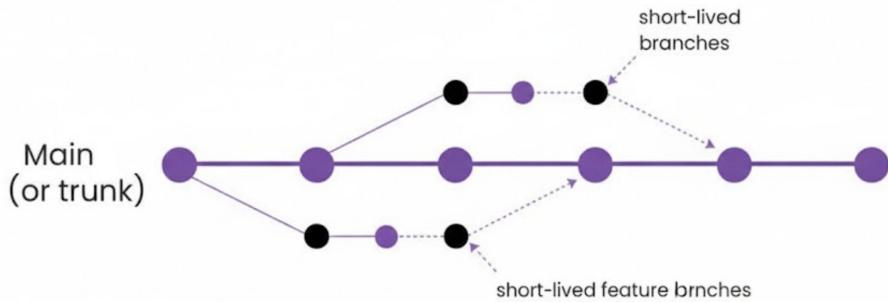
- ◆ **How it works in GitOps:** Argo CD has multiple "Applications." The Staging App tracks the `staging` branch, and the Production App tracks the `production` branch. To deploy to production, you perform a Pull Request (PR) from `staging` to `production`.
- ◆ **Focus:** Control, manual gates, and clear separation between environment states.
- ◆ **Best for:** Regulated industries or teams that require manual sign-offs before a production release.



Trunk-based is about **Continuous Deployment**, while Branch-based is about **Release Control**



# Promotions between Stages: Trunk-Based



- ◆ **No Branch Merges:** We don't merge `staging` into `production`. Both environments track `main`.
- ◆ **Environment-Specific Values:** We use separate `values.yaml`
- ◆ **Promotion = Tag Update:** Promoting a version means updating the `image.tag` in the production configuration files.
- ◆ **Decoupled:** Code is integrated into `main` immediately (CI), but deployed to production only when the configuration is updated (CD).

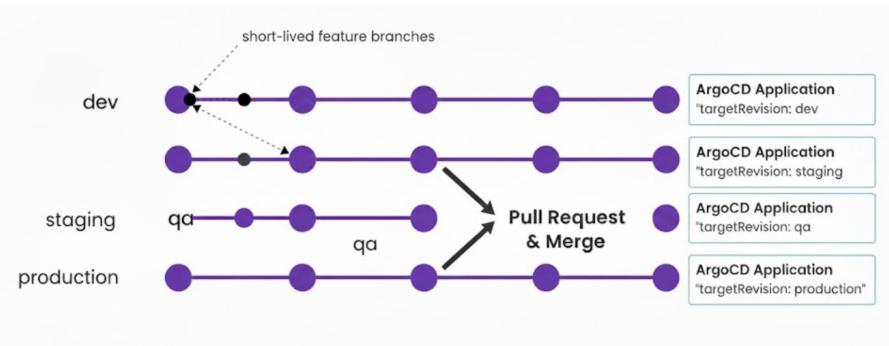
```
# staging-app.yaml (ArgoCD)
targetRevision: main
helm:
  values: |
    image:
      tag: "v1.2.0-rc1" # Promoted first for testing
  ...

# production-app.yaml (ArgoCD)
targetRevision: main
helm:
  values: |
    image:
      tag: "v1.1.0" # Remains stable until staging is verified
```

In Trunk-Based GitOps, we promote 'Artifact Versions,' not 'Git Branches.'



# Promotions between Stages: Branch-Based



```
# Application: dev-cluster
spec:
  source:
    repoURL: https://github.com/org/app
    targetRevision: dev # Watches the dev branch

---
# Application: prod-cluster
spec:
  source:
    repoURL: https://github.com/org/app
    targetRevision: main # Watches the main/prod branch
```

- ◆ **Environment = Branch:** The state of the `staging` branch is exactly what is running in the Staging cluster.
- ◆ **Promotion via Merge:** Moving code from Dev to Staging requires a Git Merge/Pull Request.
- ◆ **Visibility:** You can use `git diff dev..staging` to see exactly what is about to be deployed.
- ◆ **Permissions:** Git branch protection rules (e.g., "Only Artem can merge to `main`") act as deployment gates.

This architecture is categorized as a GitOps **anti-pattern** and is **not recommended** for production-scale environments.

# Promotions between Stages: Kargo as Bridge



- ◆ **Continuous Promotion:** Instead of manually editing YAML tags for every environment, Kargo orchestrates the "Promotion" of Freight (a bundle of Git commits, Images, and Helm charts) across stages.
- ◆ **Trunk-Based but Controlled:** It allows you to keep a single branch (Trunk-based) while providing the visual Promotion Gates (Dev → Staging → Prod) that teams usually try to solve with messy branch-based GitFlow.
- ◆ **Decoupled from CI:** Kargo stops the "CI-as-CD" anti-pattern. You don't need GitHub Actions or Jenkins to "push" changes; Kargo manages the lifecycle natively within Kubernetes.



**Before Kargo:** You write a custom bash script in GitHub Actions to `sed` an image tag in a YAML file and commit it to another branch. (Fragile & complex).



**With Kargo:** You define a Warehouse and Stages. You click "Promote" in a UI (or via CLI). Kargo handles the GitOps-compliant state change automatically.



# Argo CD: Pull-Request Generator (I)

The PR Generator is a feature of the Argo CD ApplicationSet controller.

Normally, Argo CD needs a folder in Git to create an application. The PR Generator changes this: **It treats an open Pull Request as a trigger to create a temporary Argo CD Application automatically.**

1. **The Trigger:** A developer creates a new branch (e.g., `feat-login`) and opens a **Pull Request** against the `main` branch.
2. **The Discovery:** The PR Generator constantly polls your Git provider (GitHub/GitLab). It "sees" the new PR.
3. **The Templating:** The Generator uses a template to create a new Argo CD Application on the fly. It can use variables from the PR, such as:
  - a. `\{\{ .branch \}\}`: The branch name (e.g., `feat-login`).
4. **The Deployment:** Argo CD creates a **temporary namespace** (e.g., `preview-pr-123`) and deploys the app there.





# Argo CD: Pull-Request Generator (II)

```
apiVersion: argoproj.io/v1alpha1
kind: ApplicationSet
metadata:
  name: web-app-preview-environments
  namespace: argocd
spec:
  goTemplate: true # Enables advanced templating like {{.branch_slug}}
  generators:
    - pullRequest:
        github:
          owner: my-organization      # Your GitHub Org/User
          repo: my-app-source-code    # The Repo Kargo/Argo watches for PRs
          # labels: ["preview-ready"] # Optional: Only deploy if this label is set
          requeueAfterSeconds: 60      # How often to check for new PRs
  template:
    metadata:
      # Creates a unique name like: web-app-pr-42
      name: 'web-app-pr-{{.number}}'
    spec:
      project: default
      source:
        repoURL: 'https://github.com/my-organization/k8s-manifests.git'
        targetRevision: HEAD
        path: charts/my-app
        helm:
          # This is where the MAGIC happens:
          # We override the image tag with the Git Hash from the PR
          parameters:
            - name: "image.tag"
              value: '{{.head_sha}}'
            - name: "ingress.host"
              value: 'pr-{{.number}}.preview.example.com'
      destination:
        server: https://kubernetes.default.svc
        # Deploys each PR into its own isolated namespace
        namespace: 'preview-pr-{{.number}}'
```





# Argo CD: Pull-Request Generator (III)

```
apiVersion: argoproj.io/v1alpha1
kind: ApplicationSet
metadata:
  name: web-app-preview-environments
  namespace: argocd
spec:
  goTemplate: true # Enables advanced templating like {{.branch_slug}}
  generators:
    - pullRequest:
        github:
          owner: my-organization      # Your GitHub Org/User
          repo: my-app-source-code    # The Repo Kargo/Argo watches for PRs
          # labels: ["preview-ready"] # Optional: Only deploy if this label is set
          requeueAfterSeconds: 60      # How often to check for new PRs
  template:
    metadata:
      # Creates a unique name like: web-app-pr-42
      name: 'web-app-pr-{{.number}}'
    spec:
      project: default
      source:
        repoURL: 'https://github.com/my-organization/k8s-manifests.git'
        targetRevision: HEAD
        path: charts/my-app
        helm:
          # This is where the MAGIC happens:
          # We override the image tag with the Git Hash from the PR
          parameters:
            - name: "image.tag"
              value: '{{.head_sha}}'
            - name: "ingress.host"
              value: 'pr-{{.number}}.preview.example.com'
      destination:
        server: https://kubernetes.default.svc
        # Deploys each PR into its own isolated namespace
        namespace: 'preview-pr-{{.number}}'
```





# Repository Strategies

# Repository Strategies: Mono vs Multi Repo(s)



## Mono Repo (The "One for All")

All code, Helm charts, and Kubernetes manifests for all services/teams are kept in a **single, large Git repository**.

- ◆ **How it works in GitOps:** Argo CD points to different folders within the same repo
- ◆ **Focus:** Visibility, cross-service consistency, and simplified dependency management.
- ◆ **Best for:** Smaller teams or organizations that want to share code/templates easily and maintain a global overview.

## Multi Repo (The "One per Service")

Each service or team has its **own dedicated Git repository** for its code and manifests.

- ◆ **How it works in GitOps:** Argo CD has many "Application" objects, each pointing to a different repository.
- ◆ **Focus:** Autonomy, isolation, and granular access control (RBAC).
- ◆ **Best for:** Large organizations with many independent teams where Team A should not be able to touch Team B's code or configuration.



Mono Repo prioritizes shared consistency and visibility; Multi Repo prioritizes team autonomy and security boundaries.



# Folder per Environment (I)

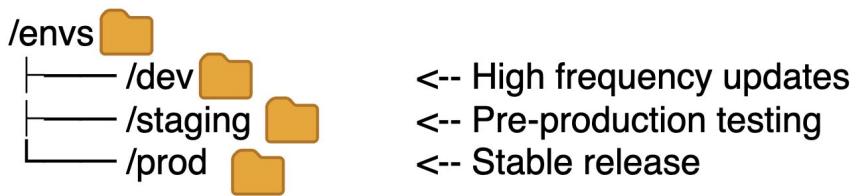
## The Core Philosophy

Instead of using Git branches (which lead to "Merge Hell"), we use folders within a single branch (usually `main`).

- **Single Source of Truth:** One branch to rule them all.
- **Promotion via Copy:** Moving a release is a simple `cp` (copy) command between folders.
- **Clarity:** A quick look at the `/envs` folder shows exactly what is deployed where.

## Simple (Stage-Based)

Best for small teams or single-region applications. It follows the software lifecycle.



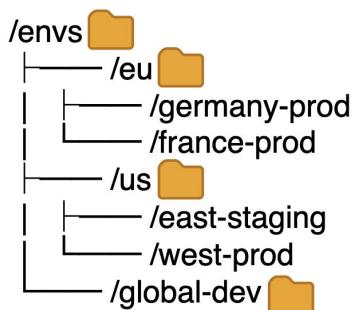
**Best for:** Projects where "Prod" is a single cluster.

**Note:** `staging` and `prod` should inherit from the same `variants/stable` folder to ensure they remain identical in configuration.

# Folder per Environment (II)

## The Geographic Matrix (Region/Country-Based)

Required for global applications with data residency (GDPR) or latency requirements.



**Best for:** Large scale apps where "Prod" is not one place, but many.

**Note:** This structure allows you to use *Kustomize Components* to apply "EU-only" settings (like specific database endpoints) to all folders under the `/eu` tree.

## The Specialized Hardware/Feature Pattern

Used when environments differ by technical capabilities (e.g., AI/ML workloads).



**Best for:** Performance testing, GPU-heavy workloads, or cost-optimization (using smaller clusters for QA).



# Folder per Environment (III)

Pattern	Complexity	Best Use Case
Stage-only	Low	Start-ups, single apps.
Stage + Region	Medium	SaaS companies with global users.
Stage + Variant	High	Specialized tech (Edge computing, AI, Finance).



# Folder per Environment (IV)

Selection Criteria	Recommended Grouping	Example Structure
Low Complexity	Stages	envs/dev, envs/prod
High Compliance	Countries	envs/germany, envs/usa
Multi-Cloud Setup	Cloud Providers	envs/aws-eu, envs/azure-us
Specific Hardware	Variants	envs/gpu-cluster, envs/cpu-cluster
Global Scaling	Regional Matrix	envs/prod-eu, envs/prod-us



# State Store

# State Store: Git, OCI or ConfigHub (I)

The State Store is the "*Source of Truth*" for your desired Kubernetes state. While Git is the default, modern GitOps architectures often use OCI or ConfigHub.

Git (The Standard)



OCI – Open Container Initiative (The Modern Way)



ConfigHub (The New Solution on the Market)



Don't confuse the **Development Source (Git)** with the **Distribution Source (State Store)**. You can write code in Git, but use a CI pipeline to push it to **OCI** for the actual deployment.

Using **OCI** as a State Store is becoming the gold standard for enterprises because it treats YAML exactly like code—versioned, packaged, and signed





# State Store: Git, OCI or ConfigHub (II)

Git



**Mechanism:** Direct sync from a Git repository (GitHub, GitLab, Bitbucket).

**Best For:** Most teams; provides full history, Pull Request workflows, etc.

**Challenge:** Large repositories with thousands of files can cause performance bottlenecks (slow polling) for the GitOps controller.

OCI



**Mechanism:** Packaging Kubernetes manifests into an OCI Artifact (the same format as Docker images) and pushing them to a Registry (GHCR, Harbor, ECR).

**Best For:** Scaling GitOps. Decouples the "Source" (Git) from the "Release" (Registry).

**Advantages:** \* **Faster Sync:** Argo CD/Flux pulls a single compressed artifact instead of cloning a whole Git history.

- **Immutable Releases:** Every version is a signed, immutable "Logical Container" of YAMLs

ConfigHub



**Mechanism:** Moves away from "File-based" Git to "Data-based" storage. It stores fully rendered, literal YAML manifests in a structured database (Source of Record) instead of managing complex templates in Git.

**Best For:** Platform Engineering teams managing hundreds of clusters or apps who want to eliminate "Config Sprawl" and template errors.

- **WYSIWYG** (What You See Is What You Get): No more hidden logic. Every manifest is stored in its final, "WET" (Write Every Time) form, making it instantly readable and validatable.

# Recap: GitOps Architecture, Patterns...



- ❑ **GitOps Architecture:** 4 principles + internal vs. external reconcilers (single cluster vs. fleet).
- ❑ **Progressive Delivery:** strategies, risks, and benefits
- ❑ **Trunk vs. Branch:** trunk reduces complexity; combine approaches when needed
- ❑ **Repository Strategies:** choose a promotion model that scales
- ❑ **State Store:** Git is common, but not the only option—consider OCI and ConfigHub



**GitOps is based on the four principles.**

But if you want to get the **most out of** it, you also need to understand the other **building block**!

Demo

