



GITOPS FOR PLATFORM ENGINEERING

# GitOps in Enterprise – Scaling and Security

MODULE 05



# Gitops at Scale – Reference Architecture(s)

# Hub & Spoke

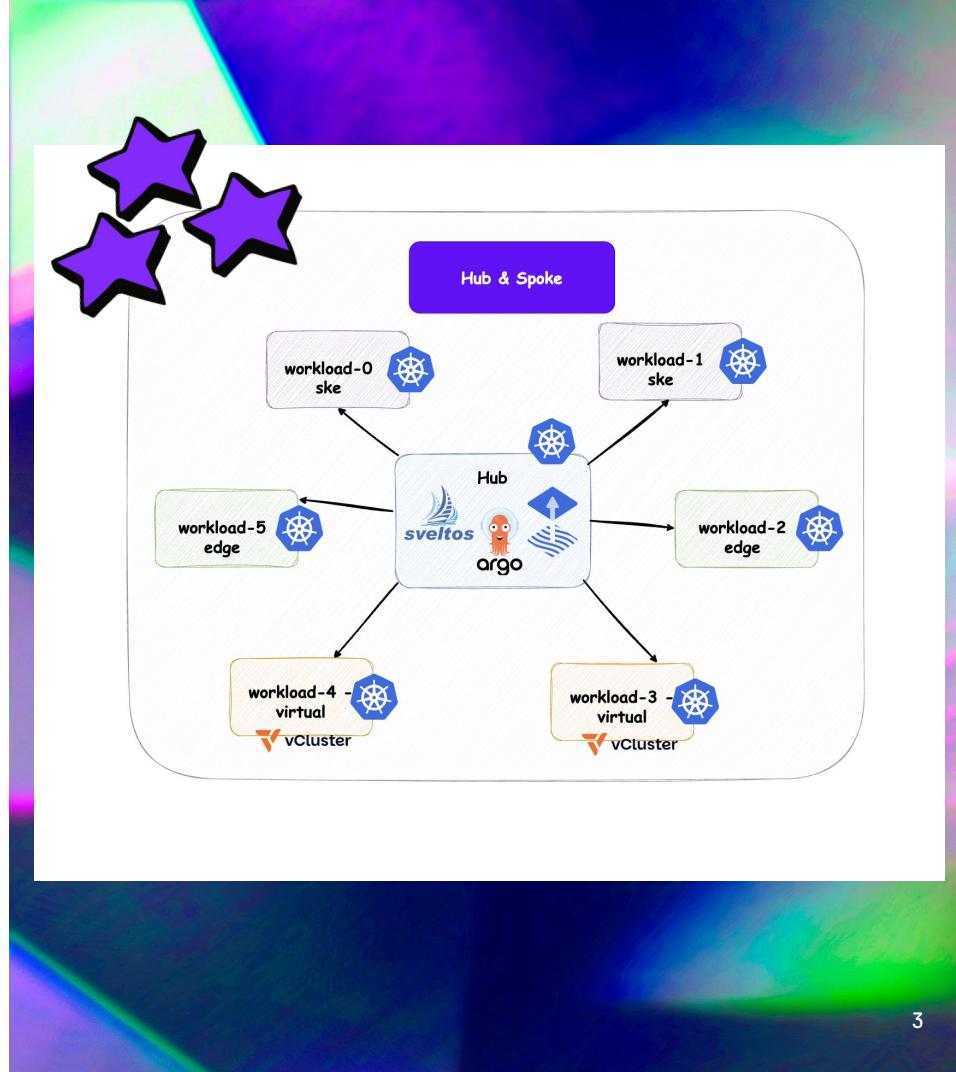
One central instance manages deployments for many remote clusters via their Kubernetes API.

## Pros:

- **Centralized Visibility:** A single "Single Pane of Glass" for all clusters and applications.
- **Ease of Management:** SSO, RBAC, and repository credentials are configured only once.
- **Efficiency:** Perfect for using **ApplicationSets** with the **cluster generator** or **ClusterProfiles** to push add-ons across the fleet.

## Cons:

- **Blast Radius:** A failure in the hub affects all connected clusters.
- **High Security Risk:** Central cluster stores admin credentials (kubeconfigs) for all target clusters.
- **Networking:** Requires the Hub to have direct network access to the target clusters' APIs.



# Instance per Cluster

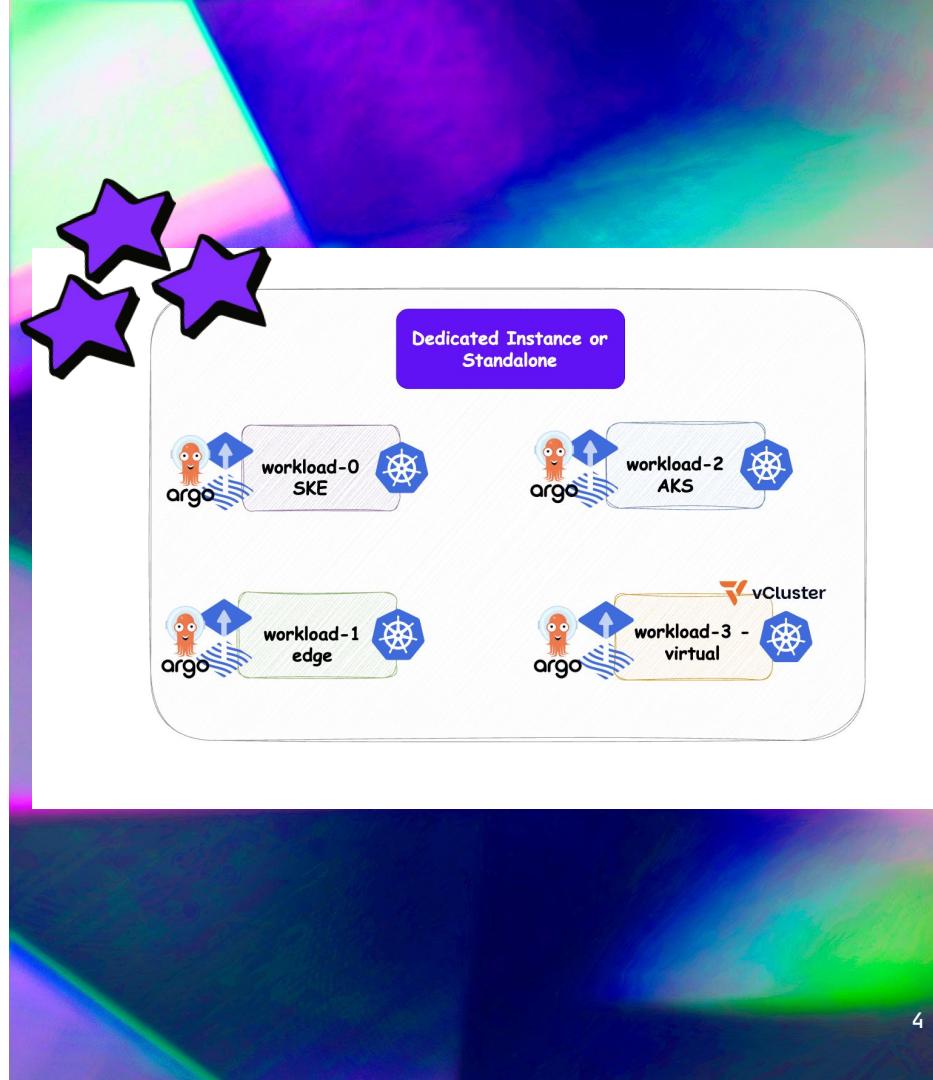
Instance is co-located with the workloads it manages.

## Pros:

- **High Reliability:** Clusters are fully autonomous. An outage in Cluster A has no impact on Cluster B.
- **Strict Isolation:** No cluster-wide admin credentials stored centrally; security boundaries match cluster boundaries.
- **Edge/Air-Gapped Ready:** Best for Edge deployments or clusters behind strict firewalls.

## Cons:

- **Management Overhead:** Every instance must be individually patched, updated, and configured.
- **Fragmented Visibility:** Developers must switch between multiple URLs/clusters to check application health.
- **Consistency Risk:** Harder to ensure all Argo/instances stay synchronized in their configuration.



# Hub & Spoke x Instance per Cluster

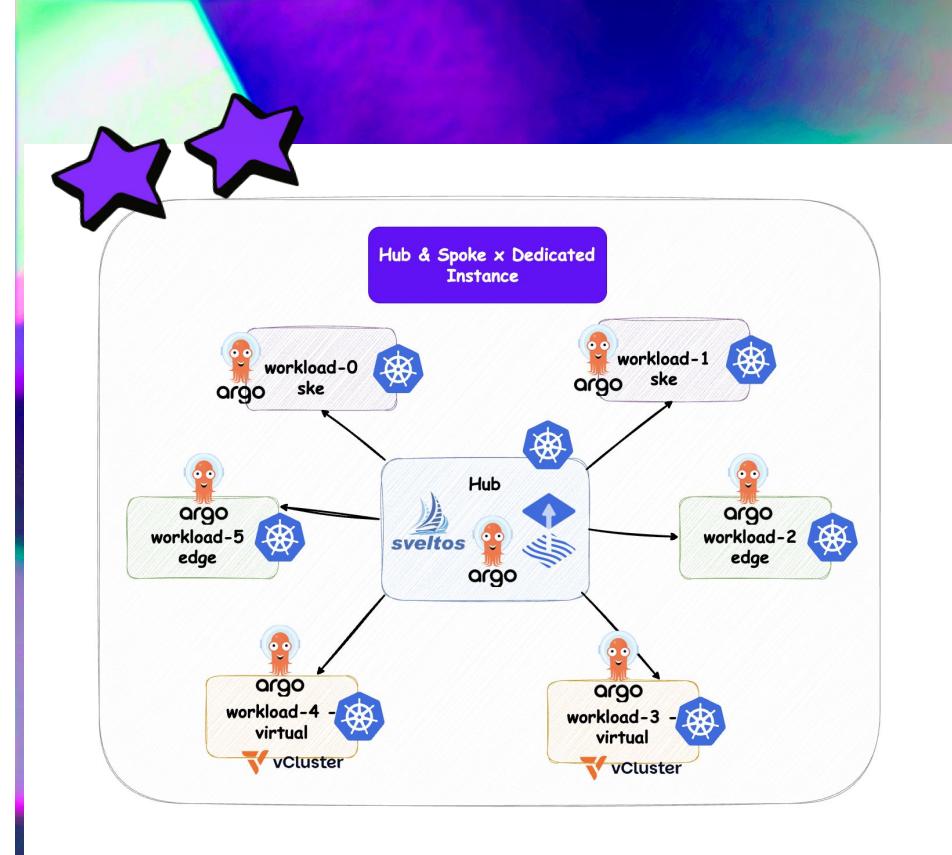
A central Hub orchestrates the fleet, while each target cluster runs its own Dedicated Instance of to execute reconciliations locally.

## Pros

- **Decentralized Reliability:** If the central Hub goes down, local instances continue to sync and maintain the desired state independently.
- **Scalability:** Heavy lifting (manifest rendering and API watching) is distributed across clusters, preventing the central Hub from becoming a performance bottleneck.

## Cons

- **Resource Overhead:** Running a full set of controllers in every single cluster significantly increases total CPU/RAM consumption.
- **Management Burden:** Operators must maintain, patch, and update two layers (the central Hub and all dedicated local instances) simultaneously.



# Hub & Spoke Agent based – Push

The Management Cluster acts as the active driver, directly "pushing" configurations to the target API servers.

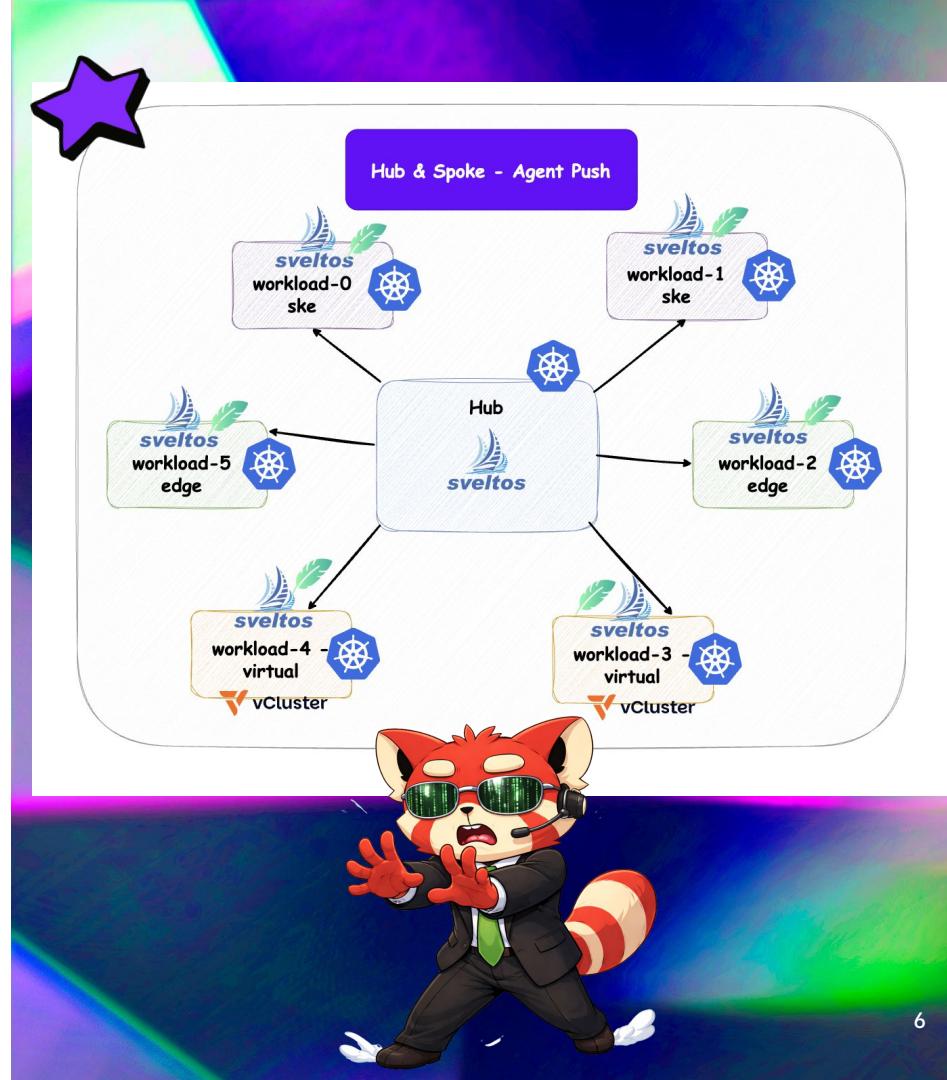
**Registration:** A **SveltosCluster** resource is created and labeled in the management cluster.

**Security:** The managed cluster's **Kubeconfig** is stored as a **Secret** within the management cluster.

**Deployment Flow:**

- Sveltos fetches the Kubeconfig Secret.
- It creates a Kubernetes client to access the managed cluster's API server.
- Resources are deployed **directly** from the hub to the spoke.

**Requirement:** The management cluster must have **direct, immediate access** to every managed cluster's API endpoint (Inbound connectivity).



# Hub & Spoke Agent based – Pull

The Managed Cluster "pulls" its own desired state from the management hub, allowing for operation behind firewalls.

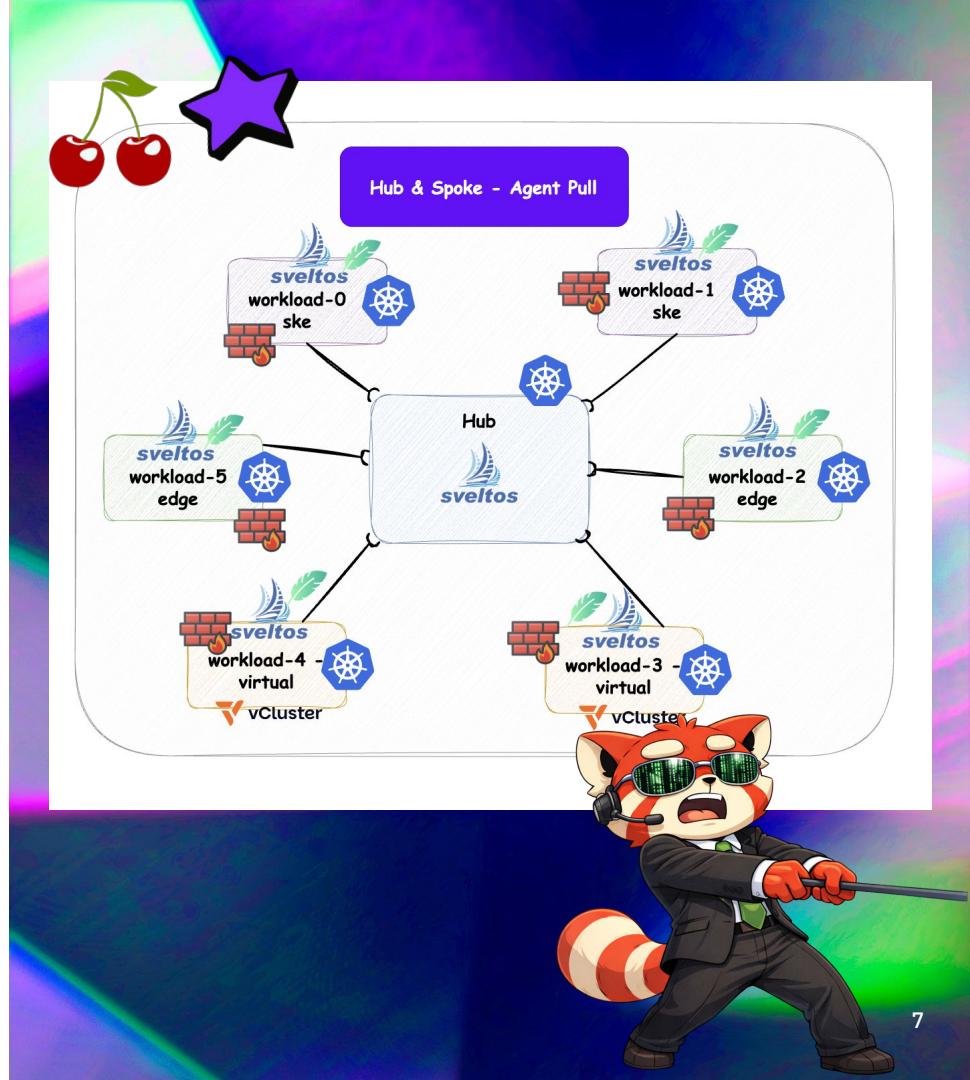
**Registration:** A **SveltosCluster** resource is created in the hub and explicitly marked for **Pull Mode**.

**Security:** A highly-permissioned **ServiceAccount** is created within the managed cluster for a local agent.

**Deployment Flow:**

- Sveltos prepares and stores resources in the management cluster.
- The **Agent** (running on the spoke) fetches the resources from the hub.
- The Agent applies the resources locally.

**Advantage:** Ideal for clusters in **private networks**; as long as the agent can reach the management hub (Outbound).





# TOO MANY CHOICES



# TOO MANY OPTIONS

# Logical Grouping

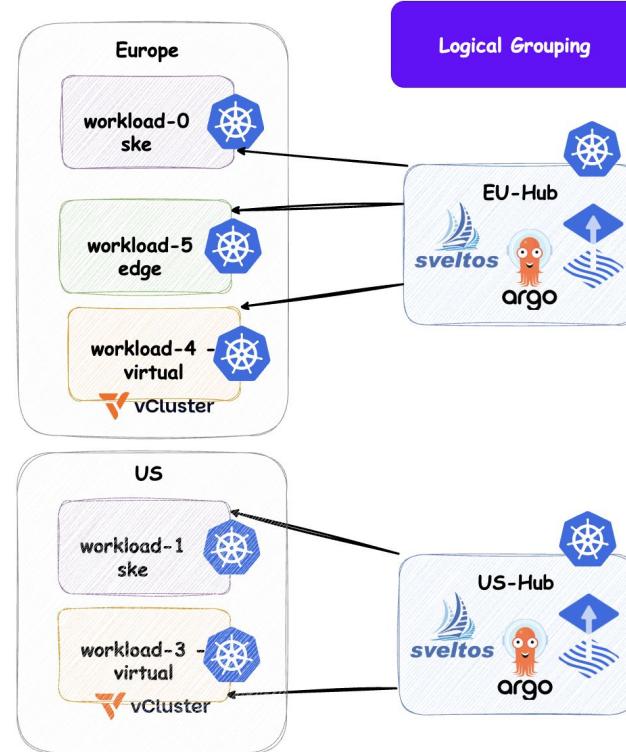
Separate instances for "Production" vs. "Non-Prod," or by geolocation units (e.g., "EU" vs. "US").

## Pros:

- **Controlled Blast Radius:** Production issues are isolated from Dev/Staging.
- **Better Scaling:** Distributes the load (Repo Server/Controller) across multiple instances.
- **Reduced Friction:** Provides a single view for specific teams while maintaining critical boundaries.

## Cons:

- **Middle Security Risk:** Central cluster(s) stores admin credentials (kubeconfigs) for all target clusters.
- **Fragmented Visibility:** No "Global Single Pane of Glass"; users must switch between different URLs/contexts depending on the group.
- **Resource Inefficiency:** Higher total CPU/RAM consumption as core controllers (API, Repo-server) are duplicated for every logical group.



# Sharding

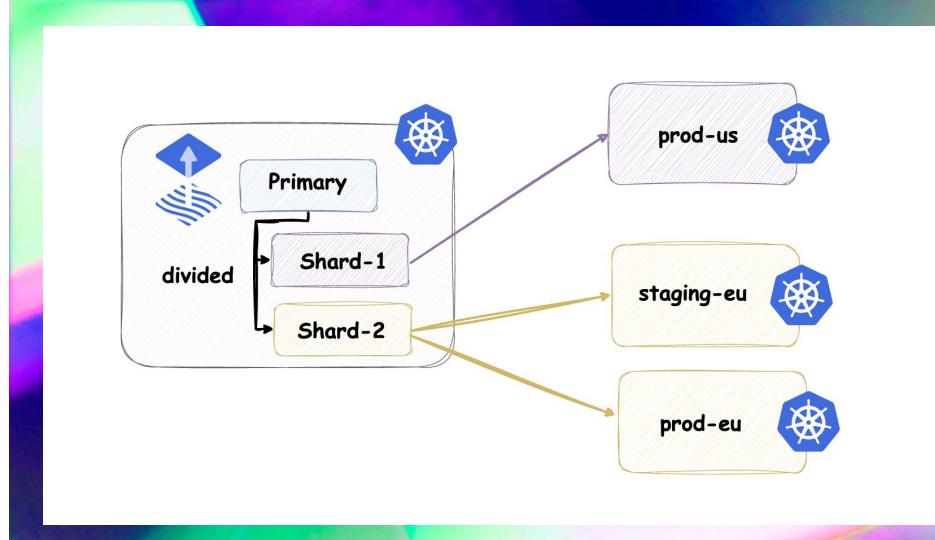
Sharding is a **horizontal scaling strategy** used to manage tens of thousands of applications by distributing the controller workload.

## Pros:

- **Massive Scalability:** Distributes CPU/RAM load across multiple controller pods to handle huge fleets.
- **Fault Isolation:** Performance issues in one shard (e.g., a massive Git repo) do not affect other shards.
- **Tenant Dedication:** Allows assigning specific shards to high-traffic tenants or critical environments.
- **High Activity:** Doesn't interfere with others.

## Cons:

- **Management Complexity:** Requires manual setup of directories and Kustomize patches for every shard.
- **Resource Overhead:** Higher total footprint as core controllers are duplicated for each shard.
- **Labeling Discipline:** All linked resources (Repo, Kustomization, Chart) must have consistent shard labels or they will be ignored.



# Sharding

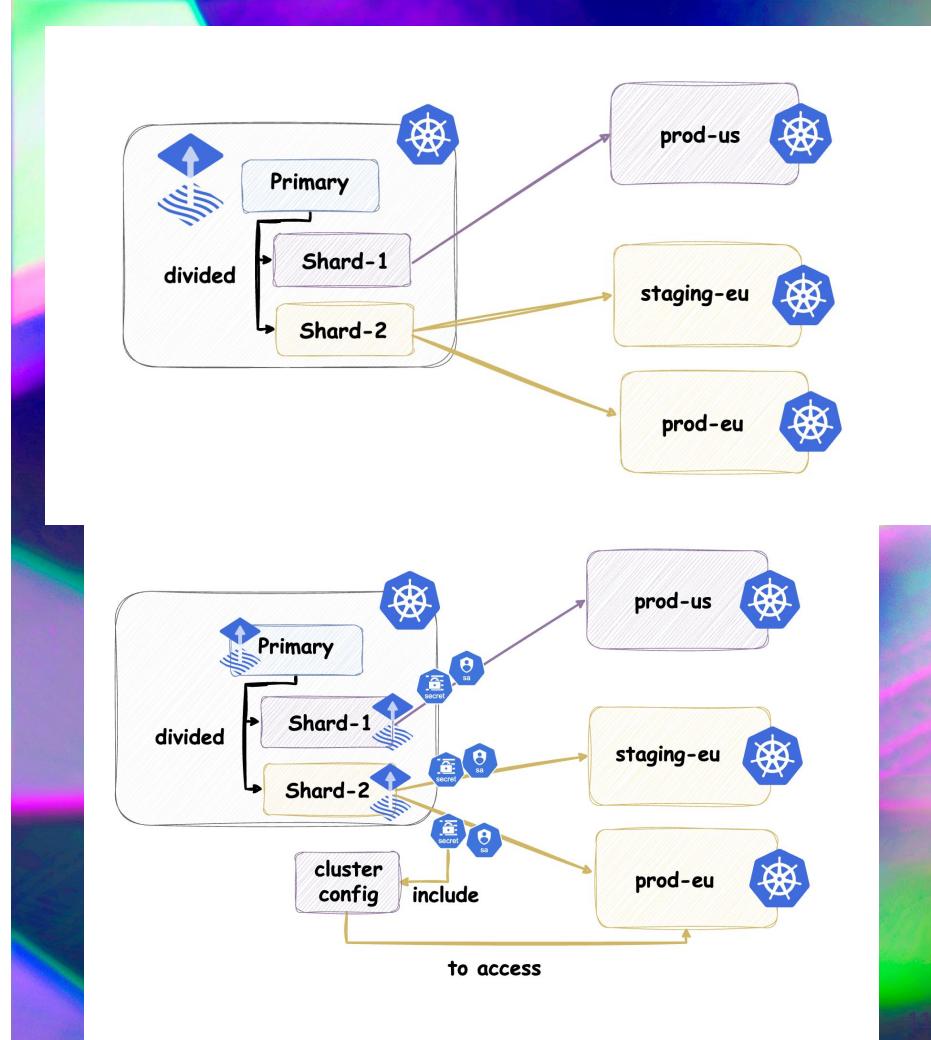
Sharding is a **horizontal scaling strategy** used to manage tens of thousands of applications by distributing the controller workload.

## Pros:

- **Massive Scalability:** Distributes CPU/RAM load across multiple controller pods to handle huge fleets.
- **Fault Isolation:** Performance issues in one shard (e.g., a massive Git repo) do not affect other shards.
- **Tenant Dedication:** Allows assigning specific shards to high-traffic tenants or critical environments.
- **High activity:** Doesn't interfere with others.

## Cons:

- **Management Complexity:** Requires manual setup of directories and Kustomize patches for every shard.
- **Resource Overhead:** Higher total footprint as core controllers are duplicated for each shard.
- **Labeling Discipline:** All linked resources (Repo, Kustomization, Chart) must have consistent shard labels or they will be ignored.





# **GitOps and Security – Secrets Management and Compliance**

# Compliance – Kyverno

Kyverno is a Kubernetes-native policy engine that manages **Compliance** and **Best Practices** using declarative resources (CRDs). It eliminates the need for complex programming languages (like Rego) by using familiar YAML-based policies.

- ◆ **Declarative Compliance:** Policies are stored in Git, versioned, and synced to the cluster just like any other application or infrastructure component.
- ◆ **Shift-Left Security:** Validates manifests during the CI/CD pipeline or at the Pull Request stage before they ever reach the cluster.
- ◆ **Automated Remediation:** Beyond just "blocking" bad configs, it can **mutate** (fix) or **generate** (create) missing resources (e.g., auto-injecting NetworkPolicies or Sidecars).



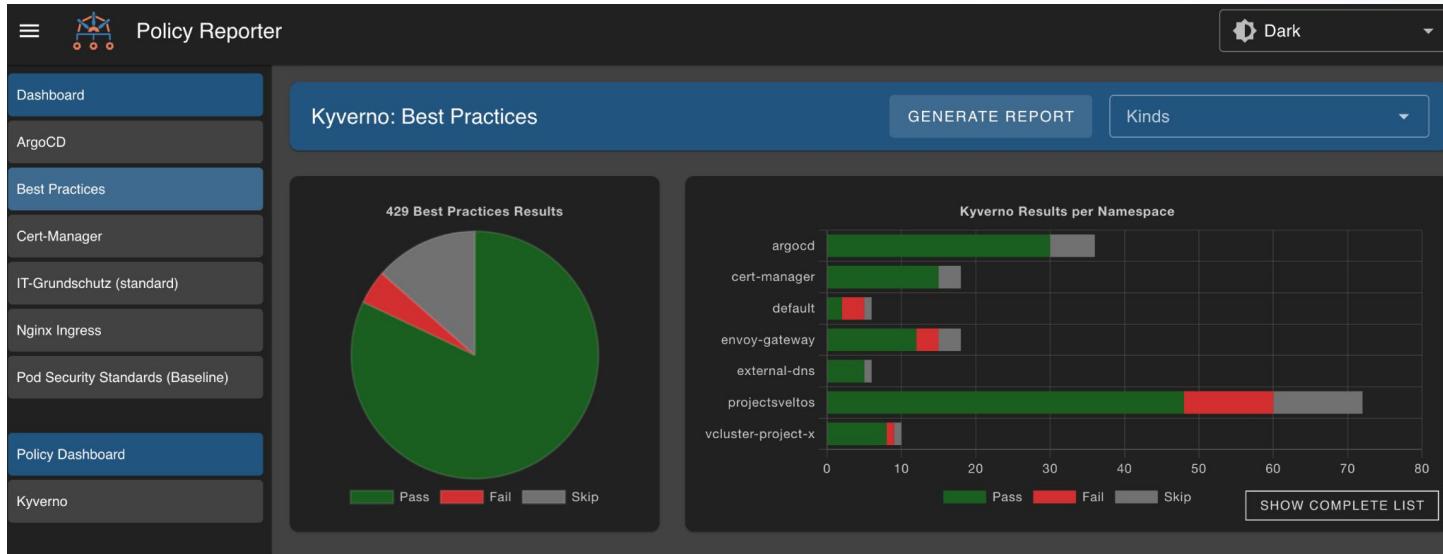
## Key Capabilities

- ◆ **Validation:** Enforces best practices (e.g., "Disallow Root User," "Require Resource Limits") by blocking non-compliant deployments.
- ◆ **Mutation:** Automatically patches incoming resources to meet organizational standards (e.g., adding mandatory labels).
- ◆ **Generation:** Creates new resources based on triggers (e.g., when a Namespace is created, generate a default Deny-All NetworkPolicy).
- ◆ **Reporting:** Provides cluster-wide compliance reports to identify existing resources that violate new security policies.



Kyverno is the 'GitOps-native Guardrail': If Git defines what *should* run, Kyverno ensures only what *is allowed* can run

# Guardrails that's Scales!



# Guardrails that's Scales! (II)

A screenshot of the Policy Reporter interface. The left sidebar includes links for Dashboard, ArgoCD, Best Practices, Cert-Manager, IT-Grundschutz (standard), Nginx Ingress, Pod Security Standards (Baseline), Policy Dashboard, and Kyverno. The main area shows a chart titled "Kyverno: Best Practices" with a progress bar at 40%. A modal window titled "Kyverno: Best Practices" is open, showing a "Catalog" section with a red button labeled "1".

Policy Reporter

Dark

Kyverno: Best Practices

GENERATE REPORT

Catalog

1

Helm Umbrella Charts

2

managed-service-catalog

3

Workload-0

Labels

Workload-N

Labels

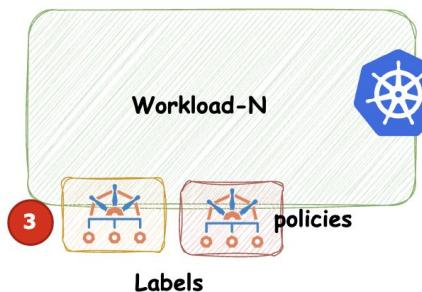
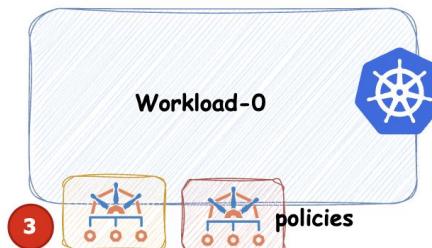


managed-service-catalog

managed-service-catalog/helm

- argo-cd
- cert-manager
- envoy-gateway
- external-dns
- external-secrets
- homer-dashboard
- kube-prometheus-stack
- kyverno
- kyverno-policies
- kyverno-policy-reporter
- loki
- template-library
- vcluster

2





# Secrets Management with GitOps

## Sealed Secrets Operator



Uses **asymmetric encryption**. You encrypt a secret locally using a public key; only the controller in the cluster holds the private key to decrypt it

### Pros:

- **Git-Centric:** The encrypted "SealedSecret" is safe to store in Git.
- **Simple Setup:** No external infrastructure (like a Vault) required.

### Cons:

- **Key Management:** If you lose the cluster's private key, you cannot decrypt your Git-stored secrets anymore.
- **Rotation:** Manual re-encryption is usually required when secrets change.

## External Secrets Operator (ESO)



Acts as an **API bridge**. It fetches the actual secret values from an external provider (AWS Secret Manager, HashiCorp Vault, Azure Key Vault) at runtime.

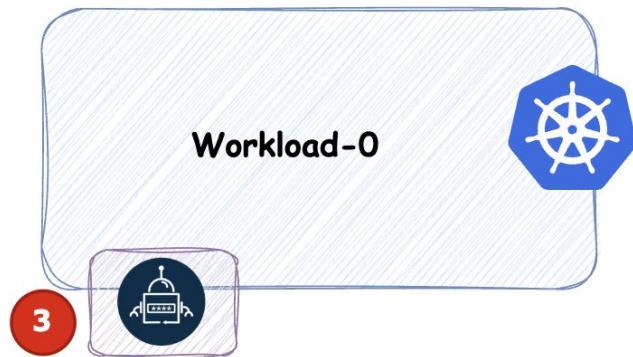
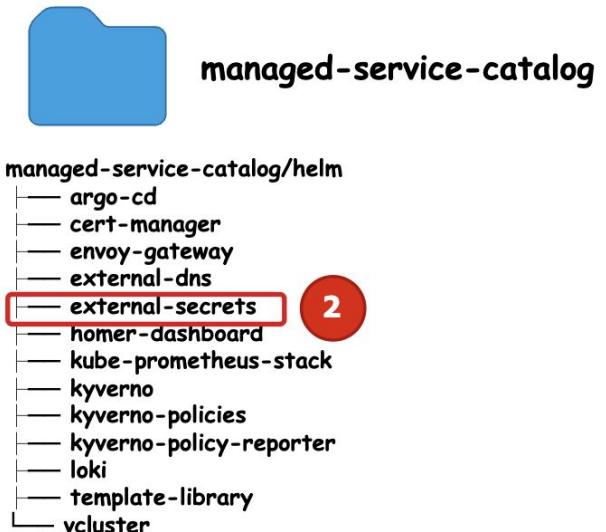
### Pros:

- **Security Best Practices:** Sensitive data never even enters your GitOps repository—only a reference (ExternalSecret) is stored.
- **Auto-Sync:** Changes in the external provider are automatically synced to the cluster.

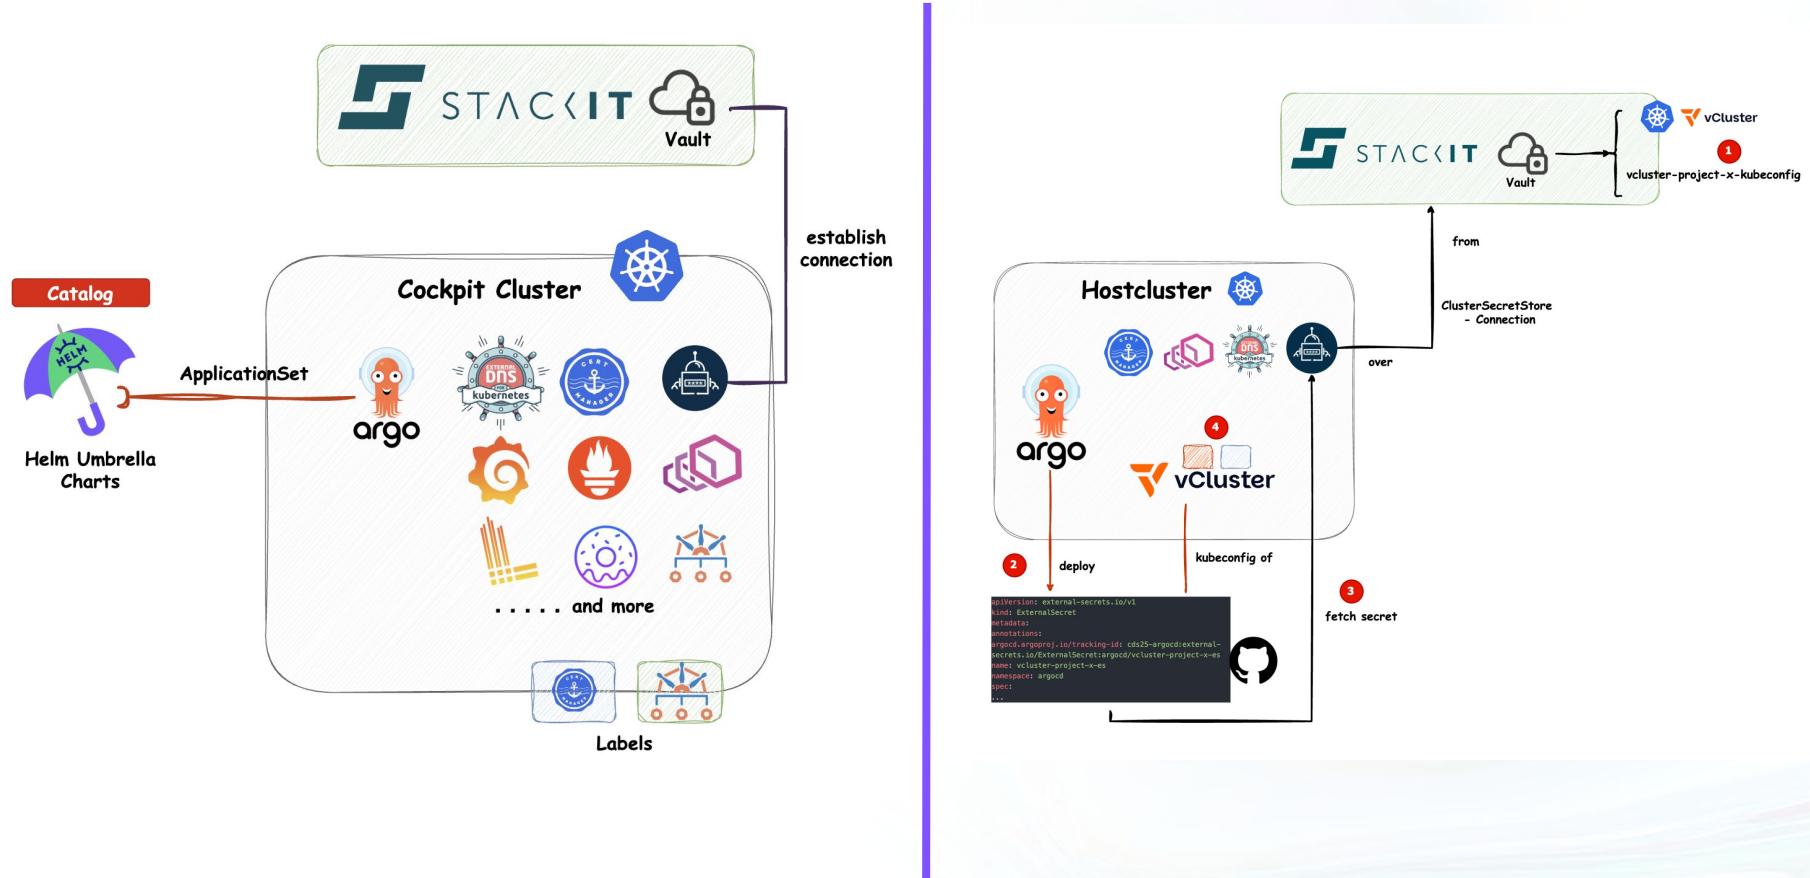
### Cons:

- **Dependency:** Requires an external Secret Store (Vault/Cloud Provider) to be available and managed.
- **Complexity:** More configuration needed (SecretStores, authentication, IAM roles).

# External Secrets Operator (ESO) – Kubara



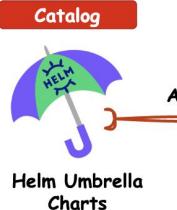
# External Secrets Operator (ESO) – Kubara



# External Secrets Operator (ESO) – Kubara



```
apiVersion: external-secrets.io/v1
kind: ExternalSecret
metadata:
  labels:
    app.kubernetes.io/part-of: external-dns
  name: external-dns-webhook-es
  namespace: external-dns
spec:
  data:
  - remoteRef:
      conversionStrategy: Default
      decodingStrategy: None
      key: dns_zone_admin
      metadataPolicy: None
      property: project_id
      secretKey: PROJECT_ID
  - remoteRef:
      conversionStrategy: Default
      decodingStrategy: None
      key: dns_zone_admin
      metadataPolicy: None
      property: sa_key_json
      secretKey: SA_KEY_JSON
```



```
secretStoreRef:
  kind: ClusterSecretStore
  name: pe-gitops-prod
target:
  creationPolicy: Owner
  deletionPolicy: Retain
  name: external-dns-webhook
```

# External Secrets Operator (ESO) – Kubara

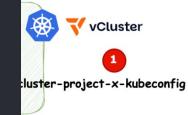


```
apiVersion: v1
data:
  PROJECT_ID: Mzg4...
  SA_KEY_JSON: eyJhY3...
kind: Secret
metadata:
  labels:
    app.kubernetes.io/part-of: external-dns
    reconcile.external-secrets.io/managed: "true"
  name: external-dns-webhook
  namespace: external-dns
type: Opaque
```

Catalog



Helm Umbrella  
Charts





# Culture Shift

# GitOps in – GitOps out



## Why GitOps Requires a Cultural Shift

- ◆ **From Manual to Declarative:** Teams must stop "fixing" things directly in the cluster (`kubectl edit`) and trust the automated process via Git.
- ◆ **Shared Responsibility:** Operations and Development merge closer; the Git repository becomes the "Single Source of Truth" for everyone, requiring high levels of collaboration.
- ◆ **Transparency & Accountability:** Every change is visible through Pull Requests, which requires a culture that embraces open peer reviews and auditability.

### Why is it hard to internalize at first?

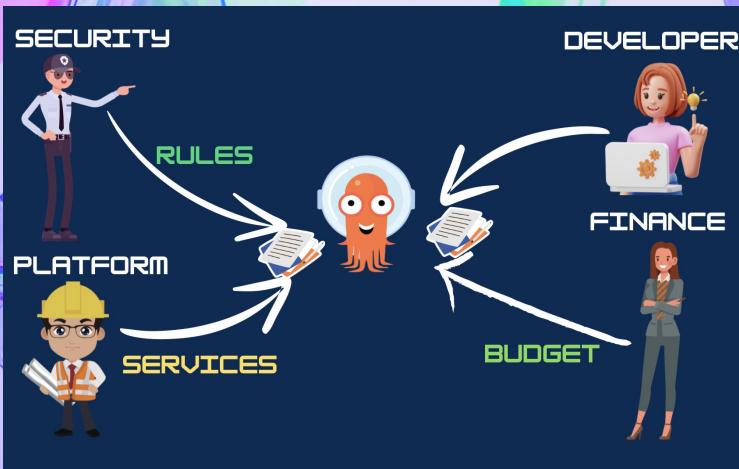
- ◆ **Relinquishing Control:** The hardest part is "letting go" of direct cluster access; developers and admins often feel restricted.
- ◆ **Steep Learning Curve:** Teams must master trunk based workflows for progressive delivery.
- ◆ **The "Lag" Factor:** Changes are no longer instant (seconds via CLI) but take time to flow through Git and the reconciler, which requires a shift in mindset regarding feedback loops.
- ◆ **Strict Discipline:** GitOps punishes "quick and dirty" fixes



GitOps is 20% Tooling and 80% Discipline. The challenge isn't the technology, but the habit of never touching the cluster manually again



# Everything as Code

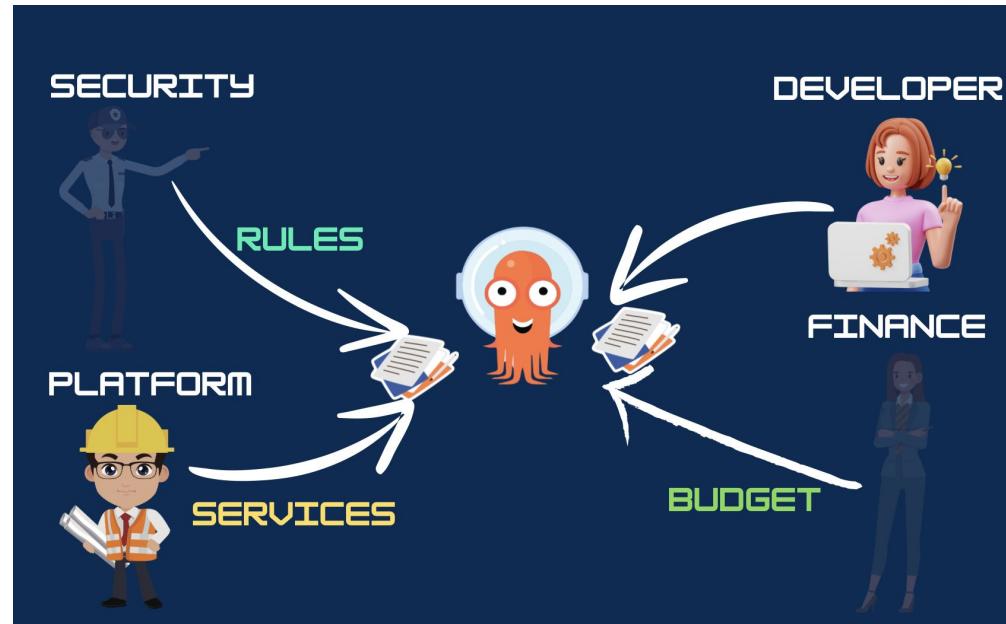


**Our vision:** A unified environment where **Security** and **FinOps** define policies or budgets as code. Whether via Git or an interface, every rule is declarative, versioned, and automatically enforced across clusters.



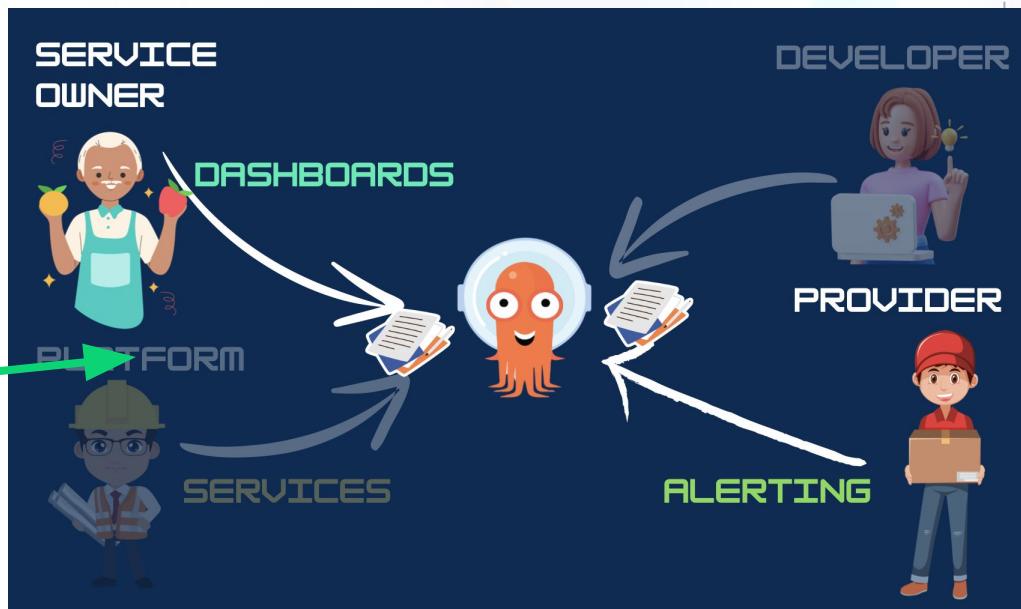
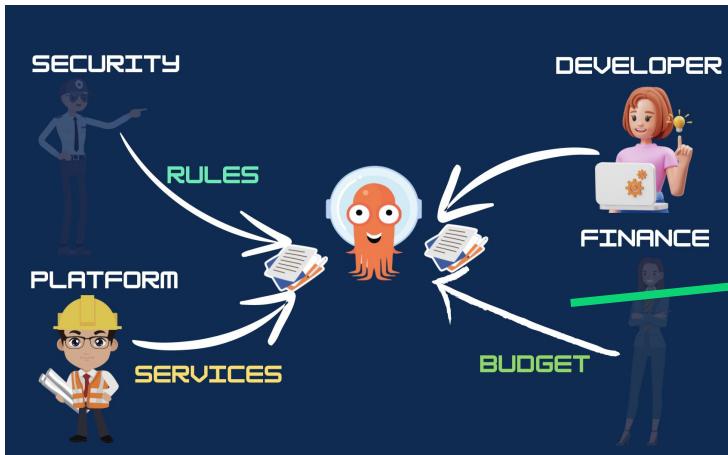


# Everything as Code (II)

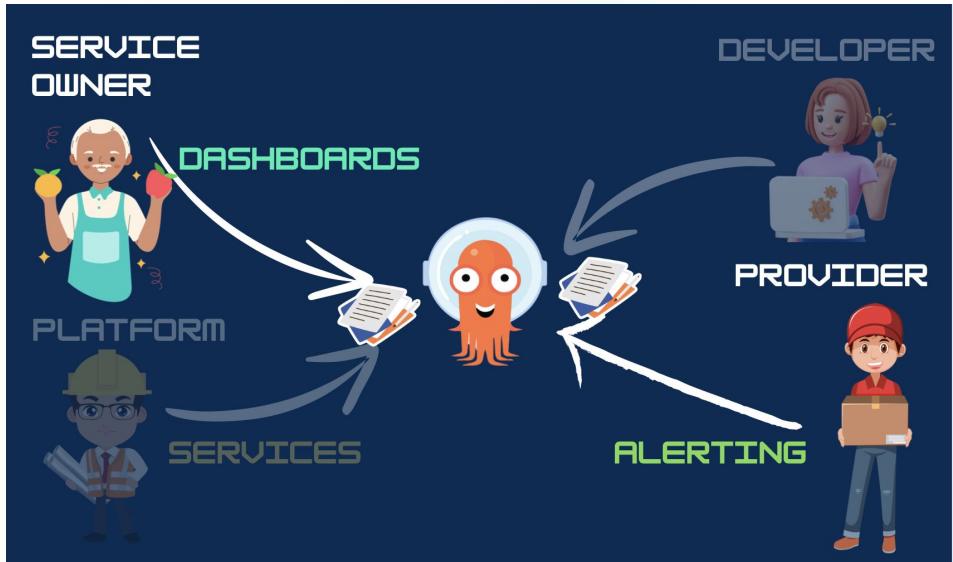




# Everything as Code (III)



# Everything as Code (IV)



Moving beyond 'ClickOps' forced a mindset shift: **Service Owners** embraced **Dashboards-as-Code**. By taking ownership through **YAML** and **Helm**, they stopped being passive users and became active creators. While the start was difficult, it ultimately transformed infrastructure into a **shared, engaging responsibility**.



# Recap: GitOps Architecture, Patterns...



- ❑ You can choose between **Hub & Spoke**, **Standalone** or ... topologies. Every environment is different, so feel free to combine these patterns to get the best of both worlds. The most popular are Hub & Spoke and Standalone.
- ❑ **Kubernetes is not secure by default.** As a platform provider, you must enforce critical policies like "*RunAsNonRoot*" through **guardrails** to keep the cluster safe.
- ❑ **Secrets management** should be as **easy as the rest of your code**. Use Sealed Secrets for a simple Git-native setup, or switch to **External Secrets Operator (ESO)** if you already have a **Vault**.
- ❑ The tools are nice, but never forget that **developers are the end-users**. No matter how you build it, the focus must stay on **providing a great self-service experience**. Culture matter.

Demo

